



703308 VO High-Performance Computing

MPI Groups, Communicators and One-Sided Communication

Philipp Gschwandtner

Overview


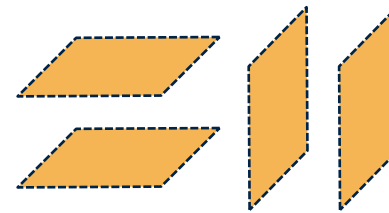

- ▶ communicators and groups
 - ▶ more ways of limiting and controlling collective communication
- ▶ one-sided communication
 - ▶ decouples data access and synchronization
- ▶ error handling

Motivation

- ▶ Real-world applications are rarely a single component
 - ▶ often MPMD
 - ▶ usually combination of libraries (e.g. molecular dynamics and quantum mechanics)
- ▶ Adds several new complexities compared to single-component software
 - ▶ collective communication via `MPI_COMM_WORLD`?
 - ▶ how to identify sub-programs?
 - ▶ how to specifically communicate between and within programs?
- ▶ And what about NUMA?
 - ▶ virtual topologies do not reflect e.g. shared memory address spaces

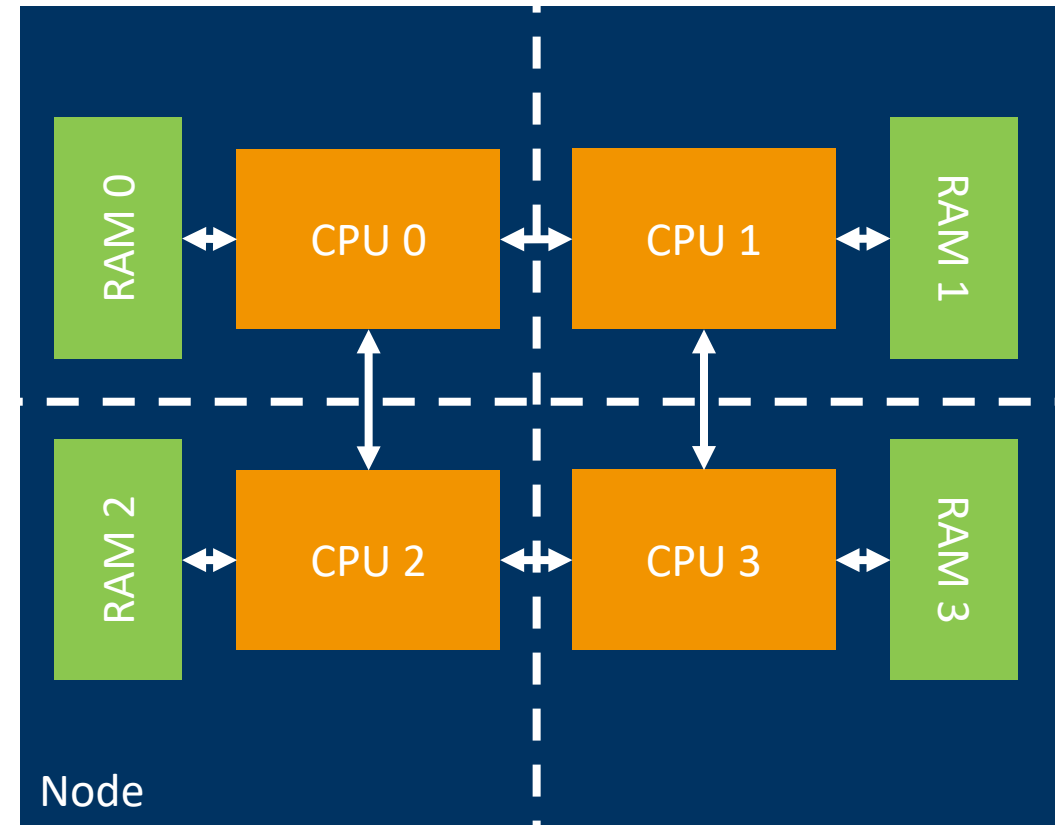
Motivation cont'd

- ▶ unstructured grid with cell and face element types
 - ▶ assume both require global communication (e.g. reduction) among their types
- ▶ a communicator per element type would be useful
 - ▶ but how?



Motivation cont'd

- ▶ consider shared memory node with 4 CPUs and many cores per CPU
- ▶ local collective communication among cores of a CPU is cheap
 - ▶ how to limit?
 - ▶ construct a communicator per CPU?
 - ▶ do this in hardware- and compiler-independent way?



Communicators and groups

- ▶ Communicators and groups hold sets of ranks
 - ▶ directly used for e.g. collective communication
 - ▶ also required for identifying single ranks
 - ▶ remember MPI basics lecture: everything in MPI is relative to a communicator or group
- ▶ Why not stick to `MPI_COMM_WORLD`?
 - ▶ isolate application sub-programs
 - ▶ individual processing steps running in parallel (task parallelism, MPMD)
 - ▶ domain decomposition (SPMD, c.f. slicing Cartesian topologies)
 - ▶ libraries (portability)
 - ▶ usability
 - ▶ add user-defined attributes such as topologies
 - ▶ performance
 - ▶ re-numbering of ranks (virtual topology vs. hardware topology)

Communicators and groups cont'd

▶ MPI_Group

- ▶ holds ordered set of ranks
- ▶ ordering is given by mapping process identifier (e.g. PID) to rank number
- ▶ construction of and operations on groups are always **local** operations

▶ MPI_Comm

- ▶ holds an MPI_Group
 - ▶ transitively holds ordered set of ranks
- ▶ can hold attributes (e.g. topology)
- ▶ constructed from groups
- ▶ construction of communicators are **non-local** operations (collectives)

Operations on MPI_Group

► Constructors

- MPI_Comm_group(...)
- MPI_Group_union(...)
- MPI_Group_intersection(...)
- MPI_Group_difference(...)
- MPI_Group_incl(...)
- MPI_Group_excl(...)
- MPI_Group_range_incl(...)
- MPI_Group_range_excl(...)

► Accessors

- MPI_Group_size(...)
- MPI_Group_rank(...)
- MPI_Group_compare(...)
 - result is MPI_IDENT, MPI_SIMILAR or MPI_UNEQUAL

► Destructor

- MPI_Group_free(...)

Operations on MPI_Comm

► Constructors

- MPI_Comm_dup(...)
- MPI_Comm_create(...)
- MPI_Comm_split(...)
- Convenience constructors such as MPI_Cart_sub(...)

► Accessors


- MPI_Comm_size(...)
- MPI_Comm_rank(...)
- MPI_Comm_compare(...)
 - result is MPI_IDENT, MPI_SIMILAR, MPI_CONGRUENT or MPI_UNEQUAL

► Destructor

- MPI_Comm_free(...)

Group and communicator workflow

- ▶ start with `MPI_COMM_WORLD`
- ▶ construct group(s) of rank subsets and modify as required
 - ▶ `MPI_Group_union()`,
`MPI_Group_range_incl()`, ...
- ▶ create new communicator from group and use for communication




Splitting communicators

- ▶ `int MPI_Comm_split(MPI_Comm comm, int color, int key, MPI_Comm* newcomm)`
 - ▶ `comm`: current communicator
 - ▶ `color`: control of subset assignment (same color: same new communicator)
 - ▶ `key`: control of rank assignment (0: sorted as in `comm`; otherwise according to ascending key values)
 - ▶ `newcomm`: new communicator
- ▶ `MPI_Comm_split_type(...)`
 - ▶ allows to split dependent on hardware properties

MPI_Comm_split example

```
MPI_Comm newComm;  
MPI_Comm_rank(MPI_COMM_WORLD, &myRank);  
int myColor = myRank / 4;  
MPI_Comm_split(MPI_COMM_WORLD, myColor,  
               myRank, &newComm);  
MPI_Comm_rank(newComm, &newRank);
```



Solutions to motivation examples

```
MPI_Comm newComm;  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
int color =  
    (elementType == TYPE_FACES);  
MPI_Comm_split(MPI_COMM_WORLD,  
    color, rank, &newComm);
```

```
MPI_Comm newComm;  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
MPI_Comm_split_type(MPI_COMM_WORLD,  
    MPI_COMM_TYPE_SHARED, rank,  
    MPI_INFO_NULL, &newComm);  
// also: OMPI_COMM_TYPE_CORE,  
// OMPI_COMM_TYPE_L1CACHE,  
// OMPI_COMM_TYPE_L2CACHE,  
// OMPI_COMM_TYPE_L3CACHE, ...
```

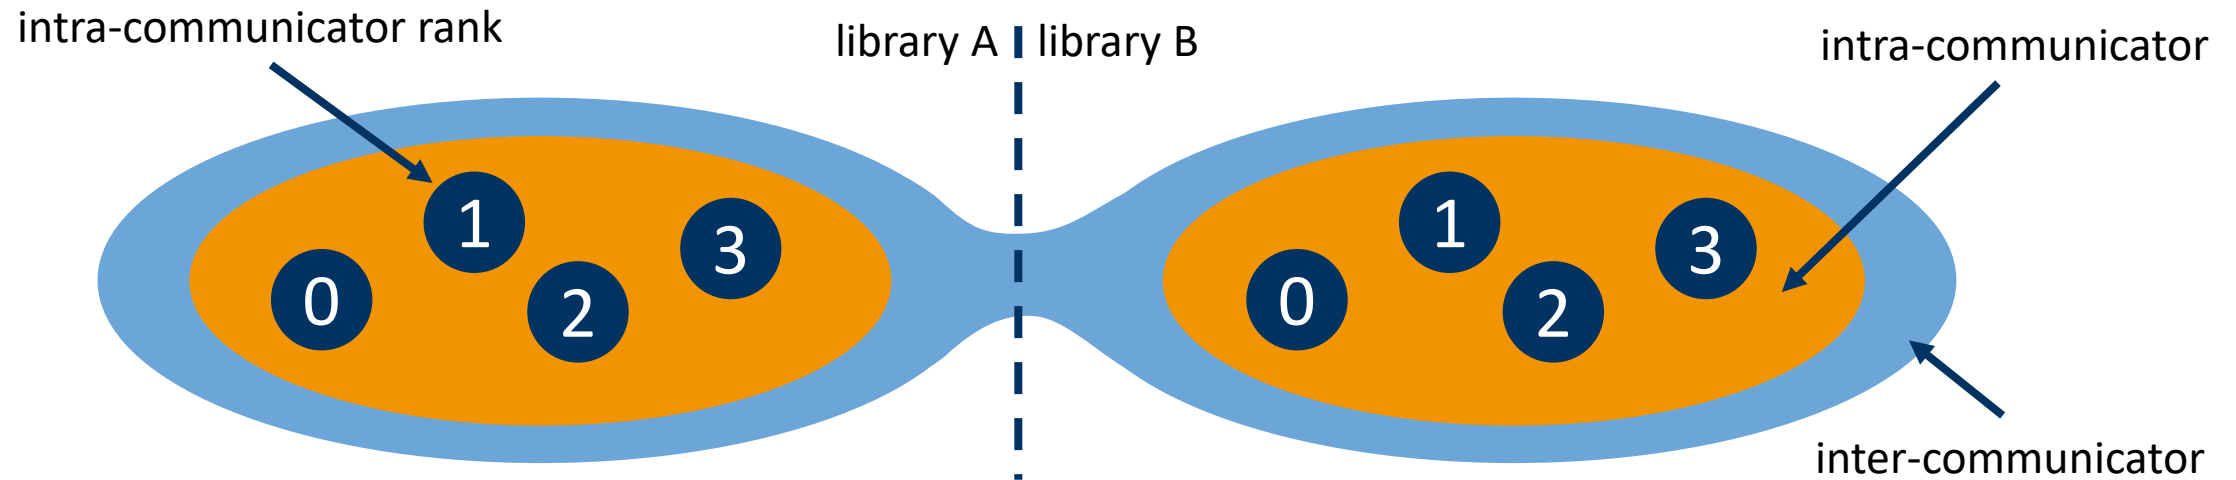
Intra- and inter-communicators

▶ intra-communicator

- ▶ collection of ranks that can send messages to each other via point-to-point and collectives
- ▶ e.g. `MPI_COMM_WORLD`

▶ inter-communicator

- ▶ collection of ranks from disjoint intra-communicators
- ▶ allows sending messages between communicators





One-sided Communication



Motivation

- ▶ **message-passing paradigm**

- ▶ fits distributed memory systems well
- ▶ data transfers among distinct address spaces require network communication
- ▶ requires explicit communication
- ▶ little control over message aggregation & synchronization

- ▶ **shared memory paradigm**


- ▶ no message passing required
- ▶ data transfer aggregation possible – write multiple bytes, elements, ... in one go
- ▶ much more convenient from a user and performance perspective
 - ▶ does not necessarily require receiving side to participate
- ▶ also, messages are needless overhead on shared-memory systems
 - ▶ send/recv function call, management, etc., just for a memcpy in the same address space?

MPI's solution: one-sided communication

- ▶ classic point-to-point (*“two-side”*) communication implies synchronization
 - ▶ at every data transfer action
 - ▶ incurs a lot of overhead in the presence of many messages
- ▶ one-sided communication decouples data movement and synchronization
 - ▶ ranks expose a *“window”* of rank-local memory
 - ▶ can be accessed by other ranks using remote memory access (RMA)
 - ▶ data accesses do not necessarily require action on the rank exposing memory
 - ▶ both read and write are possible
 - ▶ ranks no longer identify as *“sender”* and *“receiver”* but as *“origin”* and *“target”* instead


One-sided communication workflow

- ▶ allocate buffer and window
 - ▶ can ask MPI to allocate fast memory
- ▶ open window (*“start epoch”*)
 - ▶ synchronization point
 - ▶ allows data access by remote ranks
- ▶ close window (*“end epoch”*)
 - ▶ synchronization point
 - ▶ commits data accesses
- ▶ deallocate window and buffer

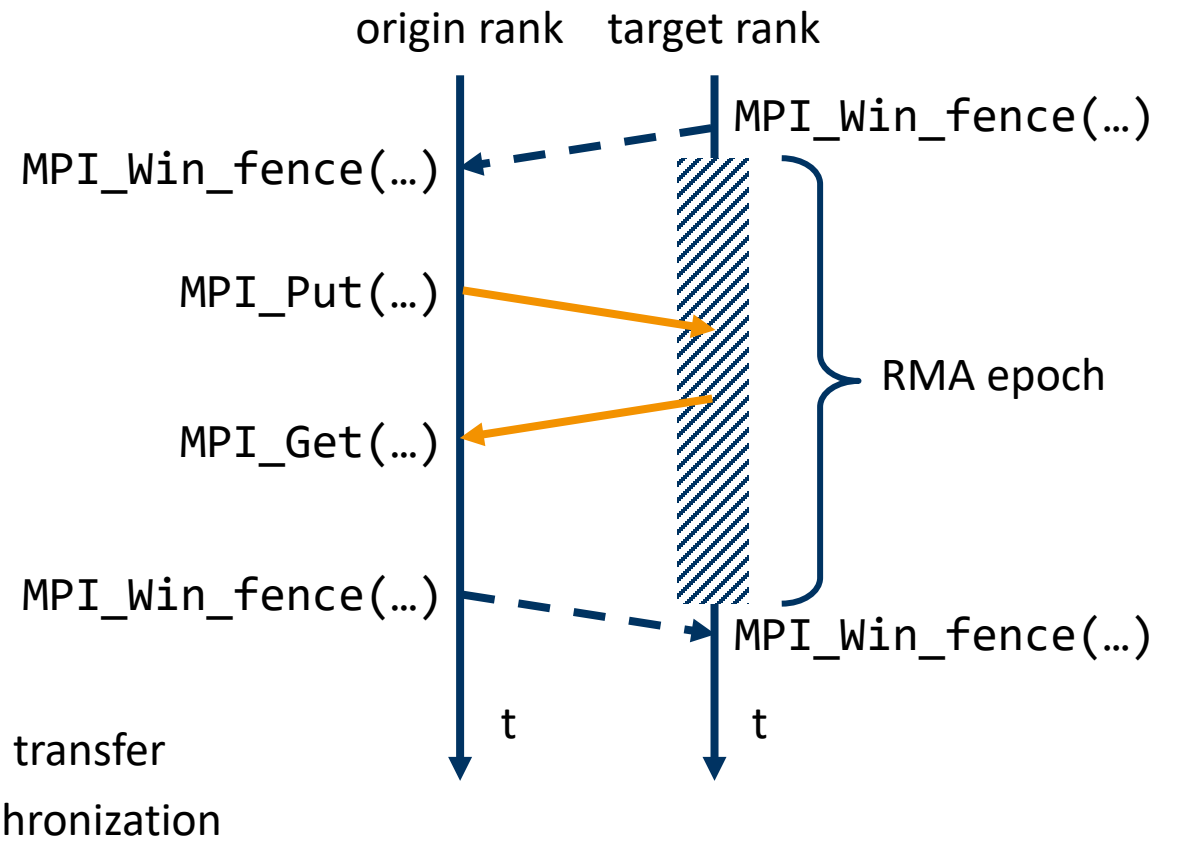


MPI one-sided communication cont'd

► two-sided



► one-sided



Means of synchronization

▶ active target synchronization

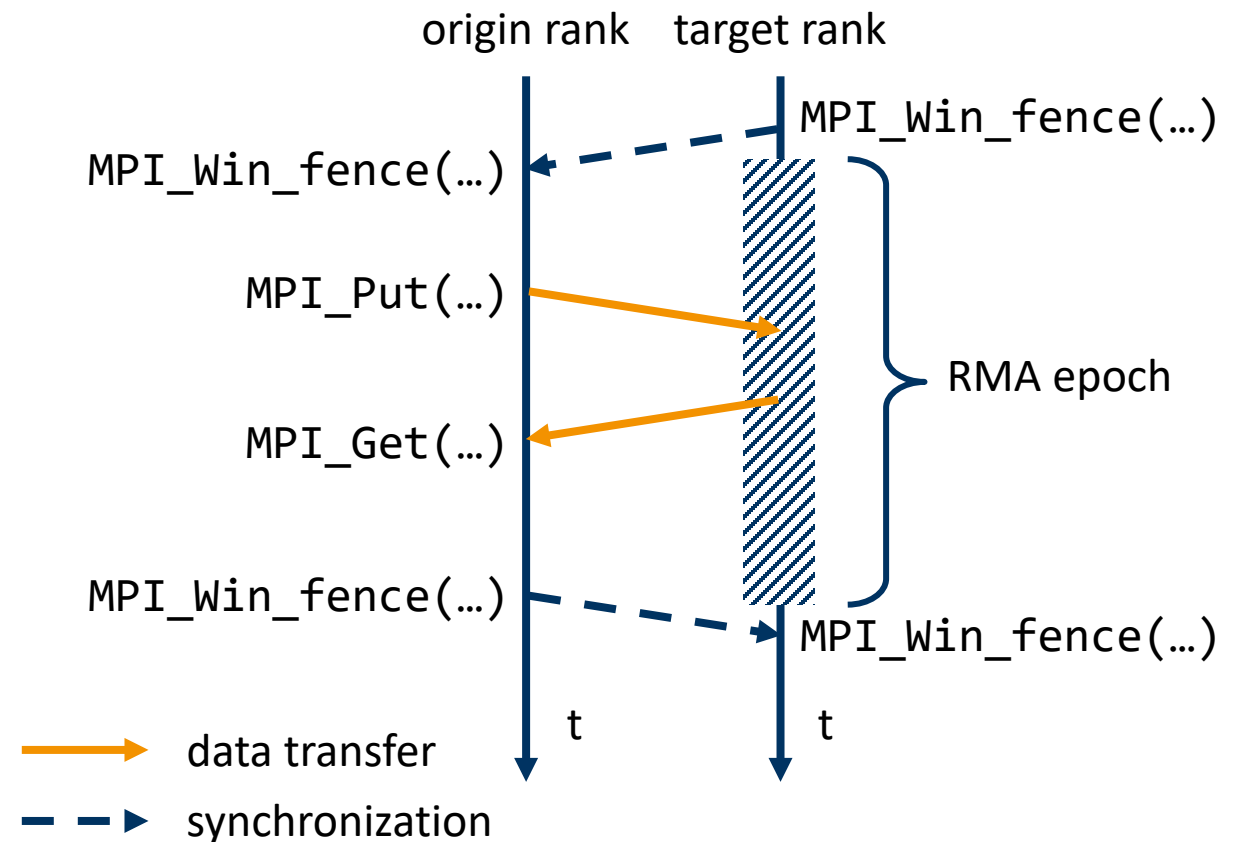
- ▶ target participates in synchronization
- ▶ similar to message-passing paradigm
- ▶ Uses either `MPI_Win_fence()` or “post-start-complete-wait”
- ▶ often preferred for bulk-synchronous parallel programs
 - ▶ all ranks execute computation and communication steps more or less in sync
 - ▶ e.g. structured grid with ghost cell exchange

▶ passive target synchronization

- ▶ target does not synchronize
- ▶ similar to shared memory paradigm
- ▶ uses `MPI_Win_lock()` and `MPI_Win_unlock()`
- ▶ often preferred for dynamic, independent access patterns
 - ▶ e.g. irregular codes

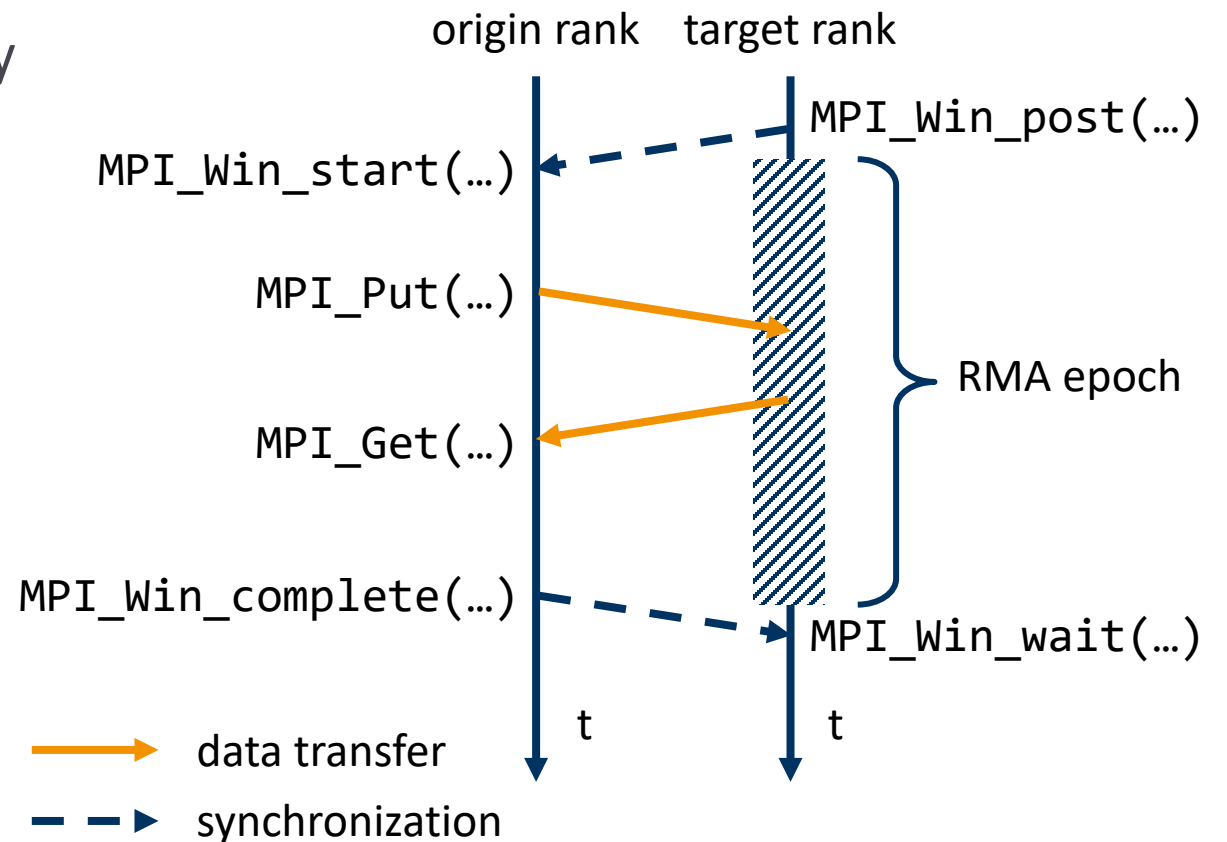
Active target synchronization: fence

- ▶ collective synchronization
 - ▶ origin/target not specified
- ▶ all control the epoch
 - ▶ starts/ends all epochs on all participating ranks
- ▶ fence enforces synchronization




Active target synchronization: post/start/complete/wait

- ▶ selective synchronization
 - ▶ origin and target specify a group they communicate with
- ▶ both control their epochs
 - ▶ origin: start/complete
 - ▶ target: post/wait
- ▶ synchronization calls may block to enforce ordering



Passive target synchronization: lock/unlock

- ▶ target neither involved in data transfer, nor in synchronization
- ▶ origin has full control over epoch
- ▶ resembles shared memory programming models (e.g. Pthreads, std::mutex, ...)
 - ▶ but not the same
 - ▶ no critical section!



Implications of one-sided communication

▶ several benefits


- ▶ allows dynamic access patterns (e.g. when target rank does not know number and ranks of origins)
- ▶ reduce synchronization overhead for multiple data transfers
- ▶ reduce management overhead on receiver side (e.g. tag matching)
- ▶ performance gain
- ▶ reduce coding effort on receiver side

▶ drawbacks

- ▶ no send/receive matching
- ▶ operations are not explicitly visible on the receiver side
- ▶ user responsible for correct order of reads/writes (race conditions)
- ▶ only non-blocking communication

Performance comparison

- ▶ LCC2, openmpi/3.1.1
2 ranks, one per node
- ▶ rank 0 sends 10^8 int to rank 1
 - ▶ once using 10^8 plain send/recv calls
 - ▶ once using 10^8 MPI_Put() calls with MPI_Fence() synchronization
- ▶ execution time reduced by 2.26x
 - ▶ only between two ranks
 - ▶ but an edge case stress test



Optional window fence assertions

- ▶ **MPI_MODE_NOSTORE**
 - ▶ local window was not updated by local store, local get or receive calls since last fence
- ▶ **MPI_MODE_NOPUT**
 - ▶ local window will not be updated by put or accumulate until next fence
- ▶ **MPI_MODE_NOPRECEDE**
 - ▶ fence does not complete any sequence of locally issued RMA calls
- ▶ **MPI_MODE_NOSUCCEED**
 - ▶ fence does not start any sequence of locally issued RMA calls
- ▶ none of these are required, but they can improve performance

Four window models

- ▶ `MPI_Win_create(...)`
 - ▶ private memory buffer already allocated, use as window
- ▶ `MPI_Win_allocate(...)`
 - ▶ allocate buffer and use as window
- ▶ `MPI_Win_create_dynamic(...)`
 - ▶ expose a buffer which is not available yet
 - ▶ use later with `MPI_Win_attach()/MPI_win_detach()`
- ▶ `MPI_Win_allocate_shared(...)`
 - ▶ allocate and use buffer in shared memory segment of the OS
 - ▶ only works for `MPI_COMM_TYPE_SHARED`

Transferring data: put & get

- ▶ `int MPI_Put(const void* origin_addr, int origin_count, MPI_Datatype origin_datatype, int target_rank, MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype, MPI_Win win)`
 - ▶ `origin_addr`: local address of data to put
 - ▶ `origin_count`: number of elements on origin side
 - ▶ `origin_datatype`: type of elements on origin side
 - ▶ `target_rank`: rank of target process
 - ▶ `target_disp`: offset of target address to base address of target window
 - ▶ `target_count`: number of elements on target side
 - ▶ `target_datatype`: type of elements on target side
 - ▶ `win`: window handle
- ▶ `MPI_Get(...)`
 - ▶ transfer data from target to origin

Transferring data: accumulate

▶ `MPI_Accumulate(...)`

- ▶ transfer data from origin and accumulate atomically at target
- ▶ can only use predefined operations of `MPI_Reduce(...)` (e.g. `MPI_SUM`)
- ▶ use `MPI_REPLACE` to get atomic put

▶ `MPI_Get_accumulate(...)`

- ▶ same as `MPI_Accumulate(...)` but store target buffer data in result buffer before accumulating
- ▶ use with `MPI_NO_OP` to implement atomic get or `MPI_REPLACE` for atomic swap

Transferring data: single-element atomics

- ▶ `MPI_Compare_and_swap(...)`
 - ▶ atomic swap if data at target buffer matches comparison value
 - ▶ must be single element
 - ▶ must be predefined integer, logical or byte type
- ▶ `MPI_Fetch_and_op(...)`
 - ▶ variant of `MPI_Get_accumulate(...)`, available for hardware optimization
 - ▶ implements a subset of `MPI_Get_accumulate(...)`'s generic functionality
 - ▶ must be single-element
 - ▶ must be predefined data type

MPI one-sided communication example

```
// rank 0 (origin)
MPI_Win window;
int buffer[SIZE] = ... ;
MPI_Win_create(&buffer, sizeof(int)*SIZE,
    sizeof(int), MPI_INFO_NULL,
    MPI_COMM_WORLD, &window);
MPI_Win_fence(MPI_MODE_NOSTORE |
    MPI_MODE_NOPUT | MPI_MODE_NOPRECEDE,
    window);
MPI_Put(&buffer, SIZE, MPI_INT, 1, 0,
    SIZE, MPI_INT, window);
MPI_Win_fence(MPI_MODE_NOSTORE |
    MPI_MODE_NOPUT | MPI_MODE_NOSUCCEED,
    window);
MPI_Win_free(&window);
```

```
// rank 1 (target)
MPI_Win window;
int buffer[SIZE] = { 0 };
MPI_Win_create(&buffer, sizeof(int)*SIZE,
    sizeof(int), MPI_INFO_NULL,
    MPI_COMM_WORLD, &window);
MPI_Win_fence(MPI_MODE_NOSTORE |
    MPI_MODE_NOPRECEDE |
    MPI_MODE_NOSUCCEED, window);
// window open, buffer can be written to
MPI_Win_fence(MPI_MODE_NOPUT |
    MPI_MODE_NOPRECEDE |
    MPI_MODE_NOSUCCEED, window);
// use buffer here
MPI_Win_free(&window);
```


RMA semantics

- ▶ order of Get and Put/Accumulate is not guaranteed
 - ▶ race condition
 - ▶ same for multiple Put operations (use Accumulate instead)
- ▶ no local access to window during access epoch
 - ▶ use an RMA operation if absolutely required
- ▶ local vs. remote completion of operation
 - ▶ no send buffer re-use after Put until end of access epoch
- ▶ no concurrent passive synchronization epochs to same target
 - ▶ only relevant in multi-threading context
- ▶ lots of MPI fence optimizations
 - ▶ where to place, which assertions to use (start with 0 and add assertions as appropriate)


Shared-memory one-sided communication

- ▶ one-sided communication also available in a truly shared memory fashion
 - ▶ origin of data transfer will get a pointer to access target memory
 - ▶ naturally only works for ranks in the same physical address space
 - ▶ c.f. POSIX shared memory segments
- ▶ allows to share memory between ranks
 - ▶ reduces memory footprint (e.g. no extra buffer for ghost cell exchange of a stencil)
 - ▶ even more efficient for intra-node communication than one-sided communication

Ghost cell exchange (message passing in shared memory)



Ghost cell exchange (MPI shared memory access)



Example for shared memory one-sided communication

```
MPI_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED,  
    0, MPI_INFO_NULL, &comm_sm);  
MPI_Win_allocate_shared((MPI_Aint)(sizeof(int) * 10),  
    sizeof(int), MPI_INFO_NULL, comm_sm, &rcv_buf_ptr,  
    &win);  
MPI_Win_fence(0, win);  
*rcv_buf_ptr = ...; // replaces the MPI_Put() call!  
MPI_Win_fence(0, win);  
MPI_Win_free(&win);
```

MPI and multithreading

- ▶ one-sided shared memory communication can become quite complex
 - ▶ there are alternatives: MPI+(OpenMP or `std::thread` or Pthreads or TBB or ...)
 - ▶ these are known as *hybrid programming models*
 - ▶ both paradigms have their ups and downs (number of programming models, compiler support, compiler optimizations, interactions, side-effects, ...)
- ▶ MPI+threads needs to be supported by MPI, indicated by one of four safety levels
 - ▶ `MPI_THREAD_SINGLE`: only a single thread per rank
 - ▶ `MPI_THREAD_FUNNELED`: multithreaded ranks, but only the main thread calls MPI
 - ▶ `MPI_THREAD_SERIALIZED`: multithreaded, but only one per time calls MPI
 - ▶ `MPI_THREAD_MULTIPLE`: multithreaded, any thread can call MPI any time (with restrictions)
 - ▶ MPI implementations are not required to support more than `MPI_THREAD_SINGLE`
 - ▶ always check beforehand with `MPI_Query_thread(...)` and call `MPI_Init_thread()` instead of `MPI_Init()`



Error Handling



Error handling

- ▶ MPI introduces additional hassle
 - ▶ imagine everything that can go wrong with a sequential process
 - ▶ add the fact that multiple processes are interacting, across the network
- ▶ default behavior: communication errors cause abort of MPI operation
 - ▶ causes respective process to exit
 - ▶ causes all other MPI processes of the same application to exit
- ▶ this only relates to halting crashes
 - ▶ also consider deadlocks or hangs

Error handling cont'd

- ▶ all MPI routines return error codes
 - ▶ `MPI_SUCCESS` if everything went well
- ▶ should always check and act accordingly
 - ▶ consider action to take when MPI calls fail
 - ▶ compare to a failing malloc in sequential program
 - ▶ make sure to free allocated resources (e.g. file handles)
 - ▶ inform the user!!

Error handling cont'd

- ▶ error behavior can be altered with `MPI_Comm_set_errhandler(...)`
 - ▶ `MPI_ERRORS_FATAL`: abort if error detected (default)
 - ▶ `MPI_ERRORS_RETURN`: return error code to user program
- ▶ user-defined error handlers can be installed
 - ▶ `MPI_Comm_create_errhandler(...)`
 - ▶ `MPI_Comm_set_errhandler(...)`
 - ▶ `MPI_Comm_get_errhandler(...)`

Additional topics

- ▶ File I/O
 - ▶ data file partitioning among processes, strided file access (derived data types!)
 - ▶ asynchronous data transfers
- ▶ process management
 - ▶ process spawning
 - ▶ socket-style communication
- ▶ Fortran-specific issues
 - ▶ slight differences in MPI function signatures
- ▶ profiling support
 - ▶ MPI_... vs. PMPI_...

Summary

- ▶ communicators and groups
 - ▶ offers high-level control over sets of ranks
- ▶ one-sided communication
 - ▶ can improve performance, but can be tricky to use
- ▶ error handling
 - ▶ graceful exits