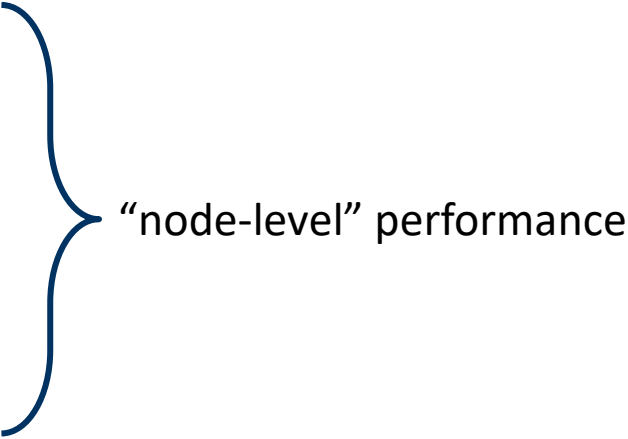




703308 VO High-Performance Computing Node-Level Performance

Philipp Gschwandtner

Motivation

- ▶ So far, we've discussed MPI, domain decomposition, load balancing, one-sided communication, etc.
 - ▶ primarily significant in inter-node performance discussion
 - ▶ What about intra-node?
 - ▶ optimizing instruction set architecture use
 - ▶ SIMD/vectorization
 - ▶ ccNUMA, data placement, cache optimizations
 - ▶ thread/core mappings
 - ▶ ...?
- 
- “node-level” performance

Overview

- ▶ On-chip heterogeneity
- ▶ Vectorization
- ▶ Common node-level pitfalls, esp. with OpenMP

On-chip heterogeneity

P core
(Golden Cove)
up to 5 GHz,
HT, private L2

E core
(Gracemont)
up to 4 GHz, no
HT, shared L2



Beginning Fall 2021

Alder Lake

Reinventing Multi Core Architecture

Single, Scalable SoC Architecture

All Client Segments – 9W to 125W – built on Intel 7 process

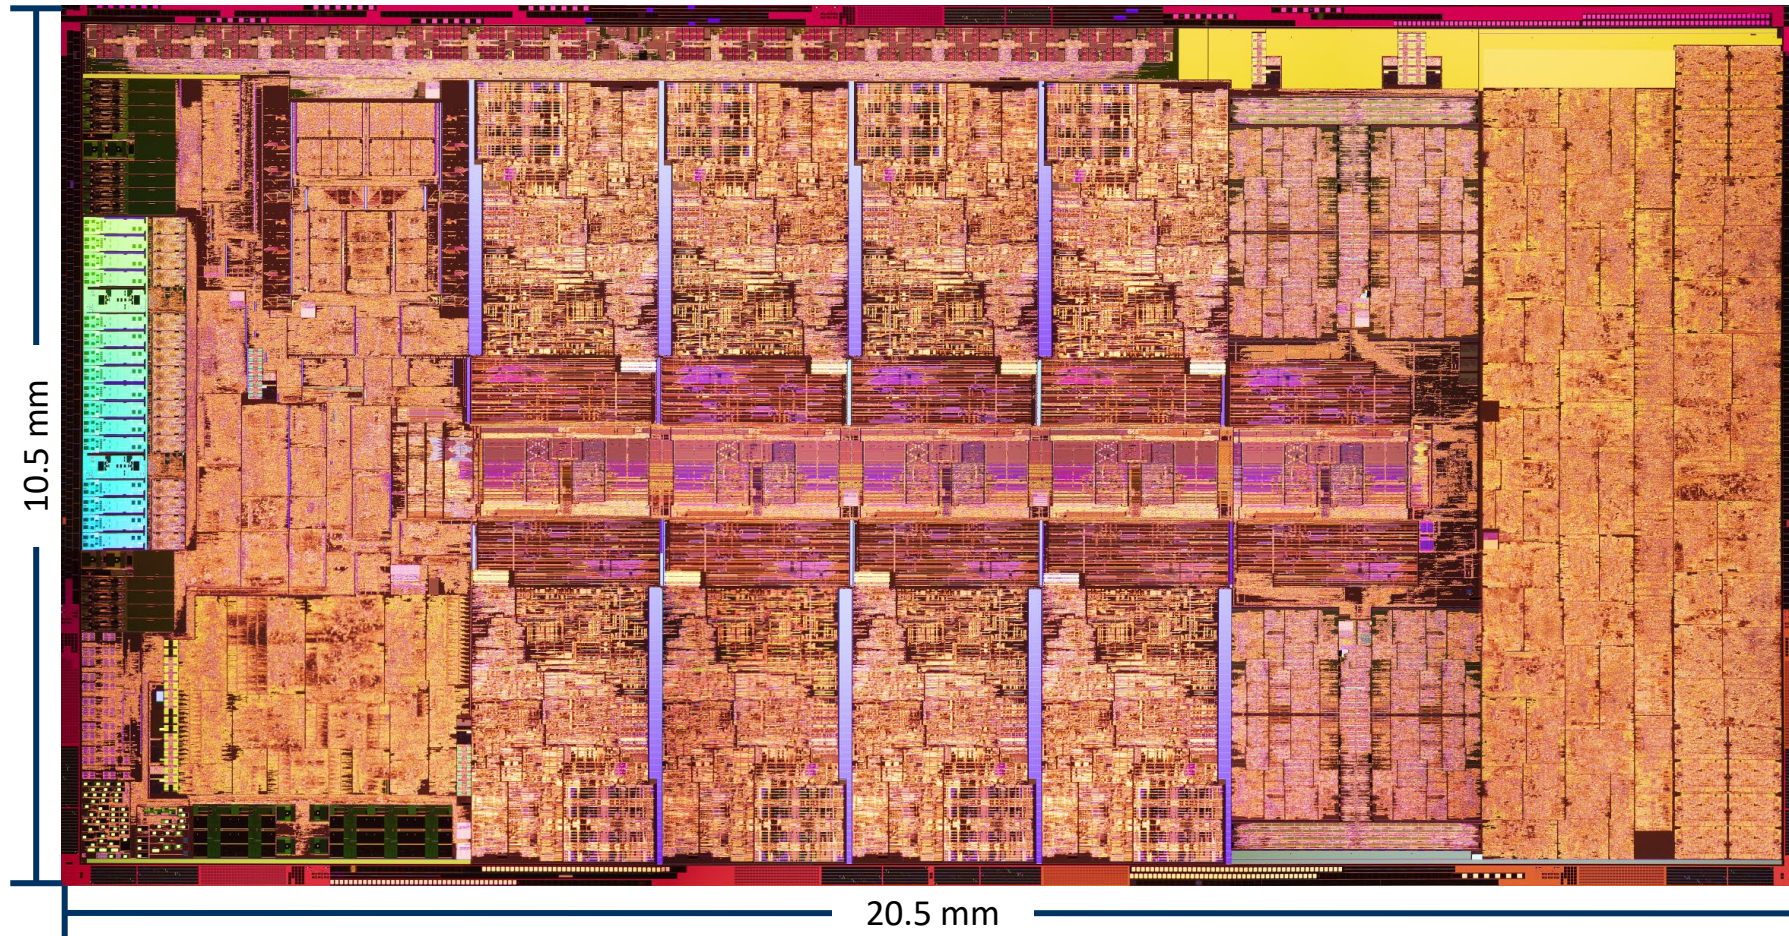
All-New Core Design

Performance Hybrid with Intel Thread Director

Industry-Leading Memory & I/O

DDR5, PCIe Gen5, Thunderbolt™ 4, Wi-Fi 6E

Alder Lake die shot



Alder Lake die shot annotated

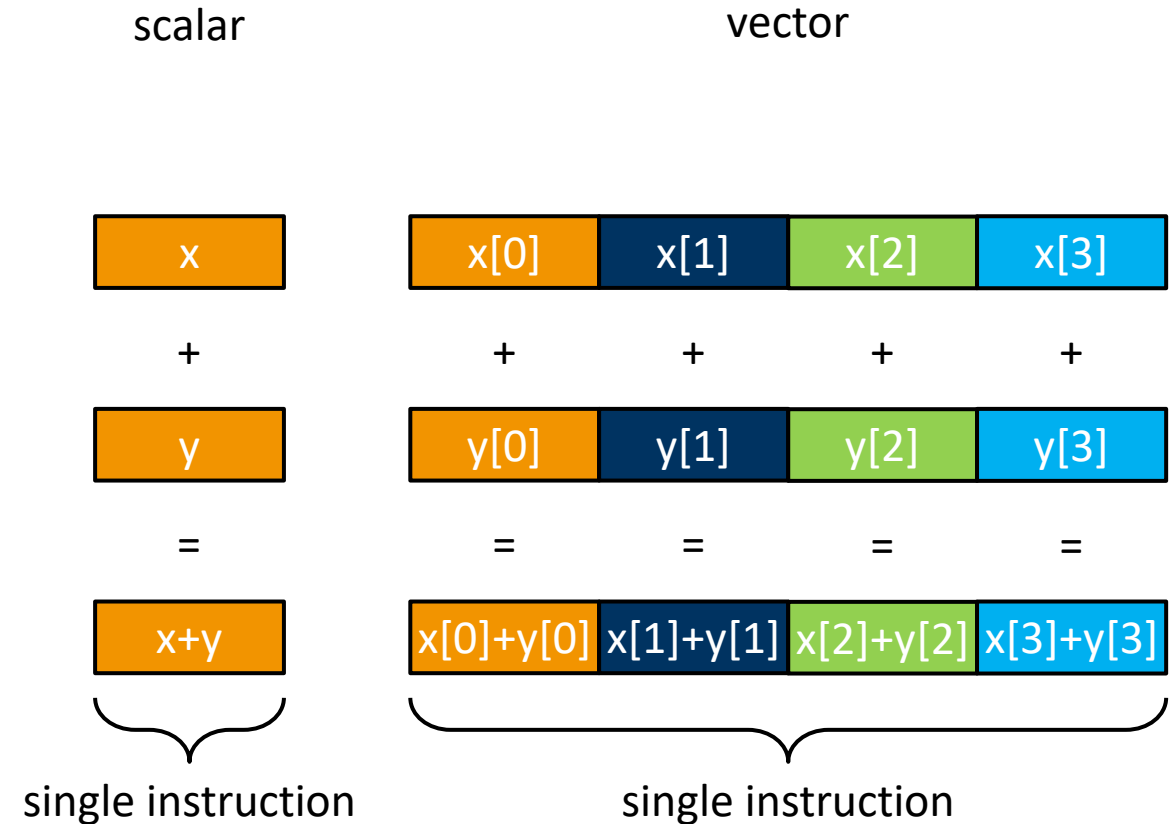


Alder Lake take-aways


- ▶ Affinity keeps gaining significance
 - ▶ fast P-cores and efficient E-cores on the same chip
(note that ARM has been doing this for years with big.LITTLE but for different reasons)
 - ▶ OS scheduler and programs need to be aware
 - ▶ e.g. previous images show an 8P+8E configuration
- ▶ L3 cache layout leads to on-chip NUMA
 - ▶ all L3 cache is accessible to every core but not with the same performance
 - ▶ has been the case at least since Haswell (~2014)
- ▶ Vectorization units take up a lot of transistor space

Vectorization

- ▶ modern CPUs have vector units
 - ▶ allow multiple data per instruction
 - ▶ performance gains of up to e.g. 4x (“SIMD width”)
 - ▶ no thread- or process parallelism involved!
- ▶ available operations and width depend on your software/hardware stack
 - ▶ hard to code manually (compiler intrinsics or assembly)




Vectorization is ubiquitous



Raspberry Pi 4 Modell B, 2GB RAM
★★★★★ 4.2 / 6 ratings
[Manufacturer link](#)

Platform	raspberry Pi
Manufacturer	raspberry Pi
Type	motherboard
SoC	Broadcom BCM2711, 4x 1.50GHz (ARM Cortex-A72)
Memory	2GB LPDDR4 RAM
Power supply	1x USB-C, 5.0V/2.5A
GPU	Broadcom VideoCore VI
Video outputs	2x Micro HDMI 2.0 (1x audio/video, 1x Video), 1x MIPI DSI
Video inputs	1x MIPI CSI2
Audio	1x 3.5mm jack (audio Out), 1x Micro HDMI 2.0
external connectors	2x USB-A 3.0 (VLI805), 2x USB-A 2.0, 1x card reader (microSDXC)



ARM SVE
Fujitsu A64FX (Fugaku's CPUs)



ARM Neon
Samsung Gear
S2 3G (~2015)



ARM Neon
Apple M1 (~2020)



Intel AVX2
Alder Lake (~2021)

Vectorization hardware cheat sheet

▶ Intel/AMD

- 64 {
 - ▶ 1997: MMX – 64 bit, integer only, 2x 32 bit or 4x 16 bit or ...
- 128 {
 - ▶ 1999: SSE – 128 bit, adds float, 4x 32 bit
 - ▶ 2000: SSE2 – 128 bit, adds double-precision (2x 64 bit) and 4x 32 bit integer
 - ▶ 2004: SSE3 – incremental update to SSE2, adds DSP and same-register
 - ▶ 2006 SSSE3 – supplement to SSE3, adds permutation and fixed-point
 - ▶ 2006: SSE4 – adds higher-level features such as dot product or population count
- 256 {
 - ▶ 2011: AVX – extension to 256 bit for floats, 3 operands
 - ▶ 2013: AVX2 – adds 256 bit integer support and fused multiply-add (FMA)
- 512 {
 - ▶ 2016: AVX-512 – extension to 512 bit, several sets of instructions depending on CPU

Vectorization hardware cheat sheet cont'd

▶ ARM

- ▶ Neon: 128 bit width, integer and single-precision float (double-precision for 64 bit CPUs)
- ▶ Helium: light-weight variant of Neon, less registers, aimed at low-power
- ▶ SVE: new & incompatible with Neon, up to 2048 bit width, directly aimed at HPC and ML
- ▶ SVE2: extends instruction set for multimedia, communication (LTE), etc. use cases

▶ Apple M1

- ▶ implements Neon

▶ IBM

- ▶ VMX/VMX128: integer and single-precision floats, up to 4x 32 bit, used in e.g. Xbox 360
- ▶ VSX: extends to 2x 64 bit and double-precision

How to check support using /proc/cpuinfo

```
$ cat /proc/cpuinfo
```

```
processor      : 0
vendor_id     : AuthenticAMD
cpu family    : 23
model         : 113
model name    : AMD Ryzen 7 3700X 8-Core Processor
stepping      : 0
microcode     : 0x8701021
cpu MHz       : 3600.000
cache size    : 32768 KB
physical id   : 0
siblings      : 16
core id       : 7
cpu cores     : 16
apicid        : 15
initial apicid : 15
fpu           : yes
fpu_exception : yes
cpuid level   : 17
wp            : yes
flags         : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2 ht syscall nx mmxext fxsr_opt pdpe1gb
rdtscp lm constant_tsc rep_good nopl tsc_reliable nonstop_tsc cpuid extd_apicid aperfmperf pni pclmuldq ssse3 fma cx16 sse4_1 sse4_2 movbe popcnt aes xsave
avx f16c rdrand hypervisor lahf_lm cmp_legacy cr8_legacy abm sse4a misalignsse 3dnowprefetch osvw wdt topoext cpb hw_pstate ibpb stibp fsgsbase bmi1 avx2 smep
bmi2 cqm rdt_a rdseed adx smap clflushopt clwb sha_ni xsaveopt xsavec xgetbv1 xsaves clzero xsaveerptr virt_ssbd overflow_recov succor smca
bogomips      : 7200.00
TLB size      : 3072 4K pages
clflush size   : 64
cache_alignment : 64
address sizes  : 48 bits physical, 48 bits virtual
power management: ts ttp hwpstate cpb eff_freq_ro
```

CPUID

- ▶ Allows you to check for CPU properties and features
 - ▶ e.g. CPU vendor, model, family, revision
 - ▶ instruction set features (e.g. AVX support)
 - ▶ other features (e.g. constant_tsc)
- ▶ virtually every CPU these days supports the CPUID instruction

```
mov eax, 0x0
```

```
cuid
```

```
; output will be:
```

```
; ebx: 0x756E6547
```

```
; edx: 0x49656E69
```

```
; ecx: 0x6C65746E
```

```
; spelling „GenuineIntel“ in ASCII
```

Controlling vectorization in software

- ▶ automatic vectorization by the compiler

- ▶ e.g. gcc: `-ftree-vectorize -fopt-info-vec-all -march=tigerlake`
- ▶ e.g. MSVC: `/Qvec-report:2 /arch:avx2`

- ▶ manual vectorization

- ▶ using OpenMP: `#pragma omp simd ...`
 - ▶ Visual Studio: requires 2019 or higher and `/openmp:experimental`
- ▶ using compiler-specific intrinsics in C/C++: `_mm_add_ps(...)`
- ▶ using ISA-specific assembly: `vaddps`

Controlling instruction sets in general

- ▶ **-march**
 - ▶ assume at minimum this architecture to be available
 - ▶ e.g. compiling `-march=tigerlake` and running on a Sandy Bridge will crash your program with SIGILL (illegal instruction), but only if modern instructions are actually emitted by compiler
- ▶ **-mtune**
 - ▶ optimize for a specific architecture but do not impose as minimum requirement
 - ▶ e.g. `-march=sandybridge -mtune=tigerlake`
- ▶ **-m<ISA-feature> or -mno-<ISA-feature>**
 - ▶ enables/disables the specified instruction set, allows more fine-grained use of ISA features compared to `-march`
 - ▶ e.g. `-mavx2`
- ▶ **Two special presets**
 - ▶ `-march=generic`: target a generic ISA of “current” CPUs (can change between compiler versions!)
 - ▶ `-march=native`: target the architecture you are currently compiling on
 - ▶ also work for `-mtune`

Verifying vectorization

- ▶ indirect means (derived metrics)

- ▶ wall time
- ▶ performance/throughput
- ▶ CPU core clock frequency (AVX)
- ▶ ...

- ▶ direct means

- ▶ compiler output (=assembly code)
- ▶ vectorization performance counters

- ▶ Be aware however, vectorized instructions exist in scalar and packed form

- ▶ packed: working on multiple data
- ▶ scalar: working on single data only

- ▶ Packed is what you're aiming for!

Reading vectorized assembly

arithmetic operation:
fused multiply-add

register order: multiply e.g. xmm1 and xmm3,
add xmm2 to result, put final result in xmm1

vfmadd132ps

AVX vector instruction

**packed variant: works
on multiple data**

data type: single-
precision floating point

vmovaps

move operation: load or store

data is aligned

What could go wrong with automatic vectorization?

- ▶ compiler-based auto-vectorization (e.g. GCC's `-ftree-vectorize`)
 - ▶ requires analysis and heuristics
 - ▶ has lots of points of failure
 - ▶ loop-carried dependencies
 - ▶ pointer aliasing
 - ▶ memory alignment
 - ▶ data type mixing
 - ▶ too complex control flow / code in general
 - ▶ numerical stability issues

```
void foo(double* a, double* b) {  
    for(int i=0; i<32; ++i) {  
        a[i] = b[i] * 2;  
    }  
}
```

OpenMP `simd` directive

- ▶ portable vectorization without compiler- or hardware-specific intrinsics
 - ▶ no need to know about GCC/LLVM/Intel/ARM...
- ▶ not the be-all end-all solution to vectorization but can help a lot
- ▶ can be combined with `for` directive
 - ▶ distributes vectorized loop iteration chunks among threads


```
// note: aligned_alloc requires -std=c11
int* a = aligned_alloc(32,
                      sizeof(int)*SIZE);
int* b = aligned_alloc(32,
                      sizeof(int)*SIZE);

// initialize a, b, and f...









#pragma omp simd aligned(a,b:32)
for(int i = 0; i < SIZE; ++i) {
    b[i] += a[i] * f;
}
```

Vectorization performance comparison

- ▶ LCC2, gcc/8.2.0, single-threaded,
 - ▶ -march=native
 - ▶ -mtune=native
 - ▶ -O2
- ▶ 10^8 vector sums on integer arrays of length 64
 - ▶ code example of previous slide
- ▶ execution time reduced by 3.84x
 - ▶ 4 integers per operation + some overhead




Vectorization hazards

-  ▶ Need data types and arithmetic operations suitable for vectorization
 - ▶ might require changing your implementation
-  ▶ Automatic only: Program analysis must deem vectorization to be safe
 - ▶ can override manually using e.g. OpenMP or intrinsics
-  ▶ Need enough instructions in stream to mitigate any static overheads
 - ▶ okay if run repetitively in a loop
-  ▶ Data should be properly aligned (though diminishing impact with each CPU generation update)
 - ▶ okay if using aligned allocations and hinting alignment to compiler
-  ▶ Code should not be memory-bound
 - ▶ okay if using loop blocking/tiling to fit data chunks in L1 cache
-  ▶ Code should not be branch-bound
 - ▶ okay if using loop unrolling (often done automatically by compiler anyway)
- ▶ Any of the above can cause vectorization to fail (not vectorized () or low performance ())

Common Node-Level (esp. OpenMP) Pitfalls

False sharing

- ▶ common performance pitfall in shared-memory programming
 - ▶ cache coherence tries to keep all data up-to-date and valid for all threads
 - ▶ can unnecessary coherence traffic and cache misses for multi-threaded programs



```
double sum = 0.0;
double sum_local[MAX_NUM_THREADS];

#pragma omp parallel
{
    int me = omp_get_thread_num();
    sum_local[me] = 0.0;

    #pragma omp for
    for (int i = 0; i < N; i++)
        sum_local[me] += x[i] * y[i];

    #pragma omp atomic
    sum += sum_local[me];
}
```

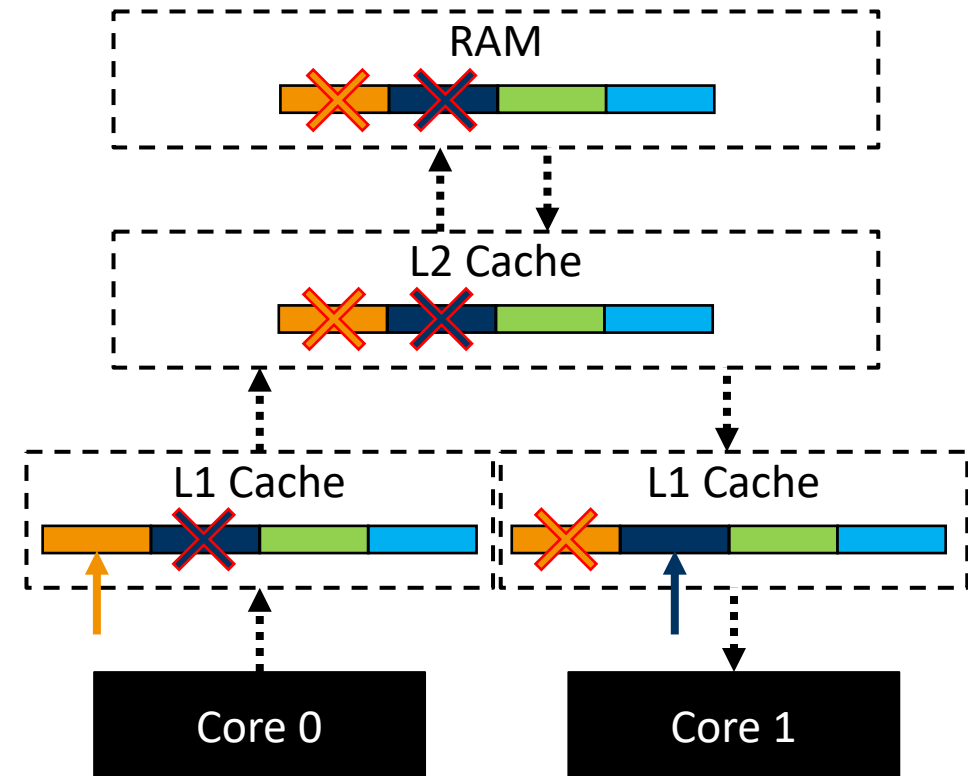
False sharing cont'd

▶ thread 1

- ▶ reads first 8 bytes
 - ▶ causes entire cache line to be fetched
- ▶ writes first 8 bytes
 - ▶ entire cache line invalidated for thread 2

▶ thread 2

- ▶ reads second 8 bytes
 - ▶ causes entire cache line to be fetched
- ▶ writes second 8 bytes
 - ▶ entire cache line invalidated for thread 1



Possible solution to false sharing: padding

```
double sum = 0.0;
double sum_local[MAX_NUM_THREADS];

#pragma omp parallel
{
    int me = omp_get_thread_num();
    sum_local[me] = 0.0;

    #pragma omp for
    for (int i = 0; i < N; i++)
        sum_local[me] += x[i] * y[i];

    #pragma omp atomic
    sum += sum_local[me];
}
```

```
double sum = 0.0;
double sum_local[MAX_NUM_THREADS][8];

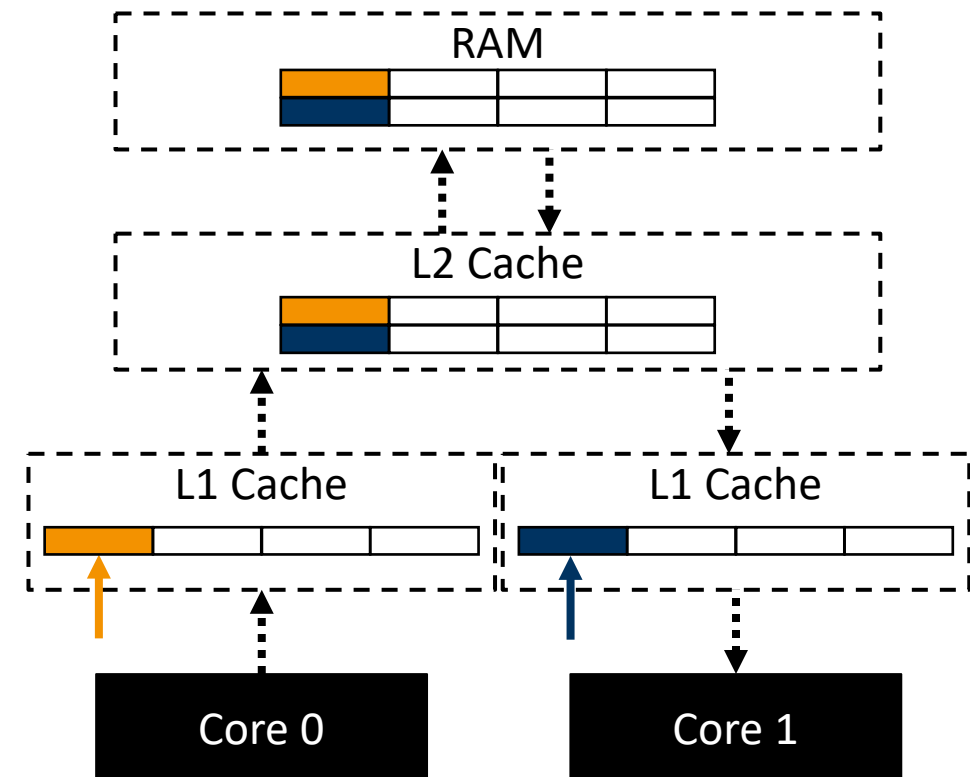
#pragma omp parallel
{
    int me = omp_get_thread_num();
    sum_local[me][0] = 0.0;

    #pragma omp for
    for (int i = 0; i < N; i++)
        sum_local[me][0] += x[i] * y[i];

    #pragma omp atomic
    sum += sum_local[me][0];
}
```


Possible solution to false sharing: padding cont'd

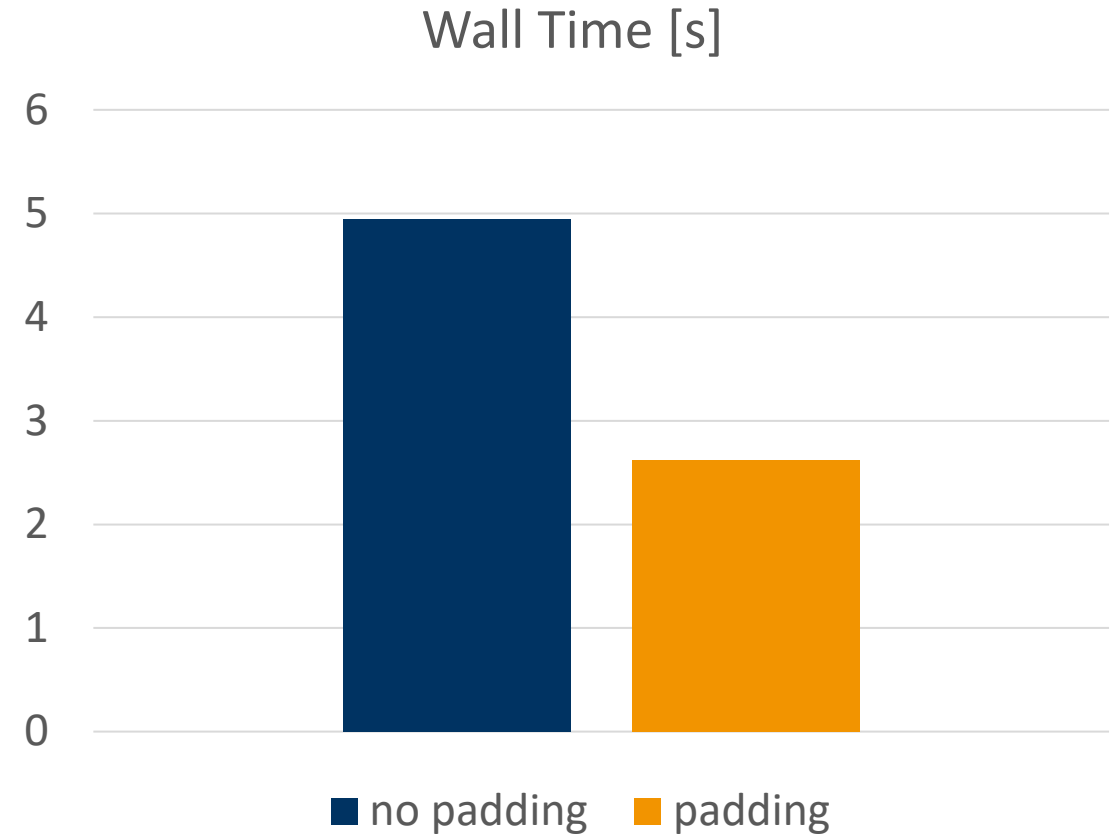
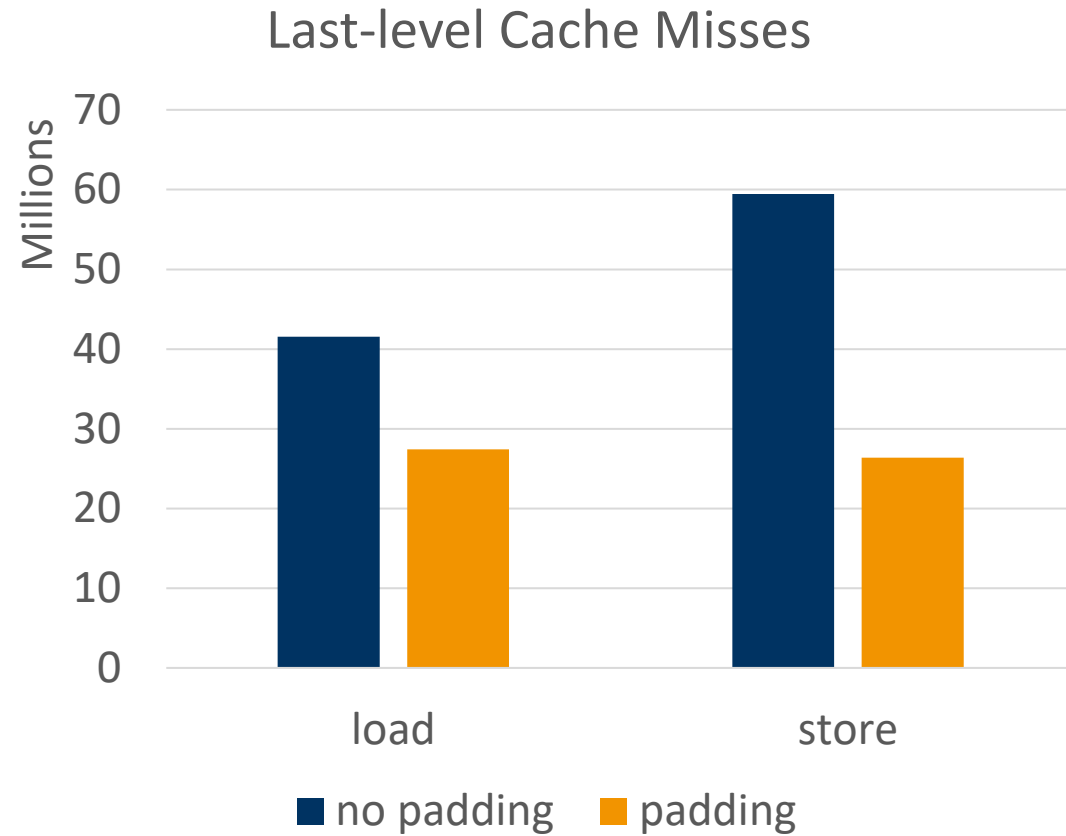
- ▶ thread 1
 - ▶ reads first 8 bytes of first cache line
 - ▶ causes first cache line to be fetched
 - ▶ writes first 8 bytes
- ▶ thread 2
 - ▶ reads first 8 bytes of second cache line
 - ▶ causes second cache line to be fetched
 - ▶ writes second 8 bytes
- ▶ drawback: memory footprint increase



How to determine padding size

- ▶ Read documentation to check for your specific CPU
 - ▶ most x86 and ARM systems use 64 byte cache lines
 - ▶ IBM POWER: 128 bytes, IBM z: 256 bytes
 - ▶ do not hard-code this, bad practice
- ▶ Better yet, ask tools, e.g.
 - ▶ `std::hardware_destructive_interference_size`: C++17 constant for L1 cache line size (constructive variant also available)
 - ▶ `cat /sys/devices/system/cpu/cpu0/cache/index0/coherency_line_size`: ask OS kernel for CPU0 and cache level 1 (index: level & instruction/data cache identifier)
 - ▶ `papi_mem_info`: executable, part of PAPI library, gives cache hierarchy information

False sharing performance comparison (LCC2, 10^9 iterations)



First touch & NUMA – how to initialize your data?

```
double* x = malloc(sizeof(double)*SIZE);
double* y = malloc(sizeof(double)*SIZE);

for(int i = 0; i < SIZE; ++i) {
    x[i] = 0.0; y[i] = 1.0;
}

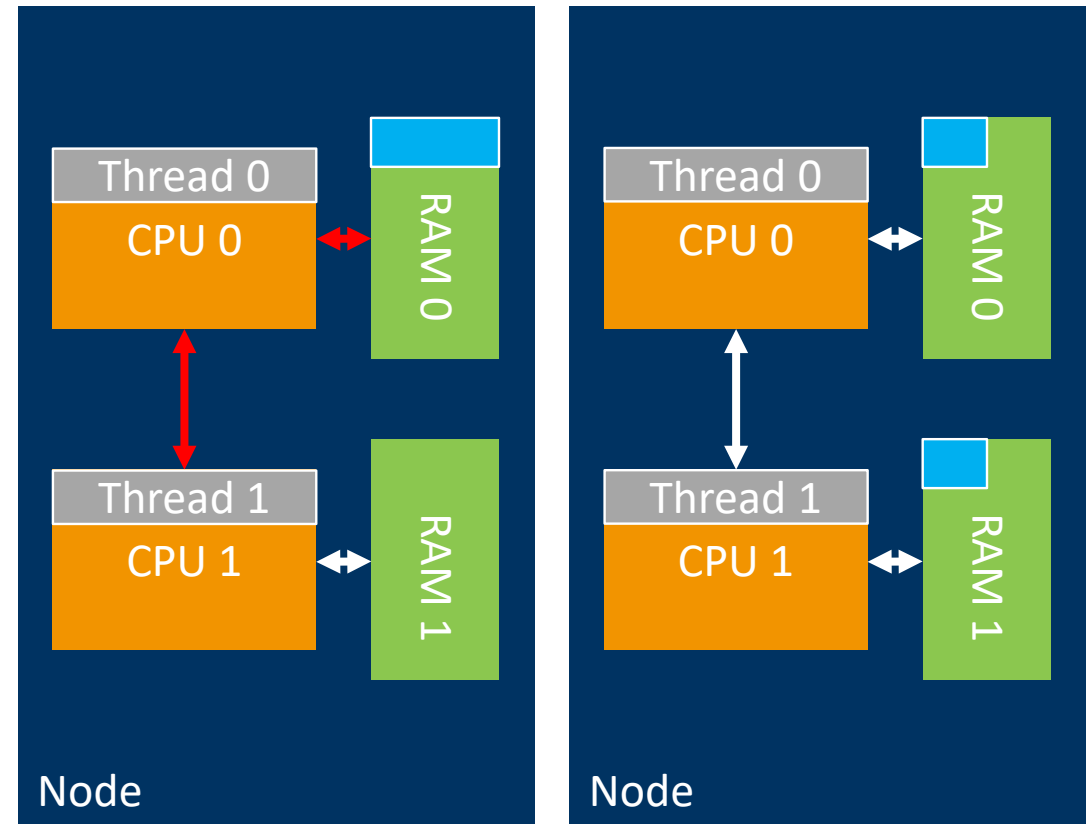
#pragma omp parallel
{
    #pragma omp for schedule(static)
    for(int i = 0; i < SIZE; ++i) {
        x[i] += y[i];
    }
}
```

```
double* x = malloc(sizeof(double)*SIZE);
double* y = malloc(sizeof(double)*SIZE);

#pragma omp parallel
{
    #pragma omp for schedule(static)
    for(int i = 0; i < SIZE; ++i) {
        x[i] = 0.0; y[i] = 1.0;
    }
    #pragma omp for schedule(static)
    for(int i = 0; i < SIZE; ++i) {
        x[i] += y[i];
    }
}
```


Sequential vs. parallel initialization on NUMA

- ▶ data is not allocated upon allocation but upon first access (*"first touch"*)
 - ▶ happens when you initialize data in the RAM module of the initializing thread
- ▶ sequential initialization
 - ▶ all data resides with RAM modules of the core of the initializing thread
 - ▶ causes bottleneck on single memory bus, additional inter-CPU traffic and higher latency for core 1
- ▶ parallel initialization
 - ▶ data resides with RAM of the threads initializing the respective chunk of data
 - ▶ only downside: need to have same domain decomposition & parallelization in initialization and computation




Performance impact of first touch and NUMA

- ▶ hudson server (2x Intel Xeon E5-2699 v3 18-core), gcc 6.3.0, 10^8 double elements, 10 repetitions
- ▶ performance improvement of compute loop (not initialization!) of 2.17x



Performance impact of first touch and NUMA cont'd

- ▶ same platform, stream memory benchmark, 3 threads per CPU
 - ▶ <https://www.cs.virginia.edu/stream/>
- ▶ between 2.08x and 2.20x higher bandwidth
- ▶ impact can vary a lot and depends also on your hardware platform



Summary

- ▶ Alder Lake, hardware architecture characteristics, implications
- ▶ Vectorization
- ▶ common shared-memory programming pitfalls

Image Sources

- ▶ Intel Architecture Day Slide: <https://download.intel.com/newsroom/2021/client-computing/intel-architecture-day-2021-presentation.pdf>
- ▶ Alder Lake Die Shots: https://www.reddit.com/r/intel/comments/qhbbow/10nm_esf_intel_7_alder_lake_die_shot/
- ▶ Raspberry Pi 4: <https://geizhals.eu/raspberry-pi-4-modell-b-a2081127.html>
- ▶ Samsung Gear S2 3G: <https://geizhals.eu/samsung-gear-s2-3g-black-a1318676.html>
- ▶ Apple M1: <https://www.computerbase.de/2020-11/apple-m1-analyse/>
- ▶ Fujitsu A64FX: <https://www.hpcwire.com/2020/02/03/fujitsu-arm64fx-supercomputer-to-be-deployed-at-nagoya-university/>
- ▶ Snail: https://live.staticflickr.com/3620/3391918403_188330c938_b.jpg